

YOGESH RAO S D

```
import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data import DataLoader
from torchvision import datasets, transforms
from torchvision.utils import make_grid
```

```
import numpy as np
import pandas as pd
from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt
%matplotlib inline
```

```
transform = transforms.ToTensor()
```

```
train_data = datasets.MNIST(root='../Data', train=True, download=True, transform=transform)
```

```
100%|██████████| 9.91M/9.91M [00:00<00:00, 11.4MB/s]
100%|██████████| 28.9k/28.9k [00:00<00:00, 340kB/s]
100%|██████████| 1.65M/1.65M [00:00<00:00, 2.72MB/s]
100%|██████████| 4.54k/4.54k [00:00<00:00, 6.69MB/s]
```

```
test_data = datasets.MNIST(root='../Data', train=False, download=True, transform=transform)
```

```
train_data
```

```
Dataset MNIST
  Number of datapoints: 60000
  Root location: ../Data
  Split: Train
  StandardTransform
  Transform: ToTensor()
```

```
test_data
```

```
Dataset MNIST
  Number of datapoints: 10000
  Root location: ../Data
  Split: Test
  StandardTransform
  Transform: ToTensor()
```

```
train_loader = DataLoader(train_data, batch_size=10, shuffle=True)
```

```
test_loader = DataLoader(test_data, batch_size=10, shuffle=False)
```

```
class ConvolutionalNetwork(nn.Module):
```

```
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(1,6,3,1)
        self.conv2 = nn.Conv2d(6,16,3,1)
        self.fc1 = nn.Linear(5*5*16,120)
        self.fc2 = nn.Linear(120,84)
        self.fc3 = nn.Linear(84,10)
```

```
    def forward(self, X):
        X = F.relu(self.conv1(X))
        X = F.max_pool2d(X, 2, 2)
        X = F.relu(self.conv2(X))
        X = F.max_pool2d(X, 2, 2)
        X = X.view(-1, 5*5*16)
        X = F.relu(self.fc1(X))
        X = F.relu(self.fc2(X))
        X = self.fc3(X)
        return F.log_softmax(X, dim=1)
```

```
torch.manual_seed(42)
```

```
model = ConvolutionalNetwork()
```

```
model
```

```
ConvolutionalNetwork(
  (conv1): Conv2d(1, 6, kernel_size=(3, 3), stride=(1, 1))
  (conv2): Conv2d(6, 16, kernel_size=(3, 3), stride=(1, 1))
  (fc1): Linear(in_features=400, out_features=120, bias=True)
  (fc2): Linear(in_features=120, out_features=84, bias=True)
```

```

    (fc3): Linear(in_features=84, out_features=10, bias=True)
)

for param in model.parameters():
    print(param.numel())

54
6
864
16
48000
120
10080
84
840
10

criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr=0.001)

import time
start_time = time.time()

# Variables ( Trackers)
epochs = 5
train_losses = []
test_losses = []
train_correct = []
test_correct = []

# for loop epochs
for i in range(epochs):

    trn_corr = 0
    tst_corr = 0

    # Run the training batches
    for b, (X_train, y_train) in enumerate(train_loader):
        b+=1

        # Apply the model
        y_pred = model(X_train) # we not flatten X-train here
        loss = criterion(y_pred, y_train)

        predicted = torch.max(y_pred.data, 1)[1]
        batch_corr = (predicted == y_train).sum() # Trure 1 / False 0 sum()
        trn_corr += batch_corr

        # Update parameters
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        # Print interim results
        if b%600 == 0:
            print(f'epoch: {i}  batch: {b} loss: {loss.item()}')

    # Detach the loss tensor before appending to the list
    train_losses.append(loss.detach())
    train_correct.append(trn_corr)

    # Run the testing batches
    with torch.no_grad():
        for b, (X_test, y_test) in enumerate(test_loader):

            # Apply the model
            y_val = model(X_test)

            # Tally the number of correct predictions
            predicted = torch.max(y_val.data, 1)[1]
            tst_corr += (predicted == y_test).sum()

    # Calculate the loss for the test batch within the no_grad context
    test_loss = criterion(y_val, y_test)
    # Detach the test_loss tensor before appending to the list
    test_losses.append(test_loss.detach())
    test_correct.append(tst_corr)


current_time = time.time()

```

```
total = current_time - start_time
print(f'Training took {total/60} minutes')
```

```
epoch: 0 batch: 600 loss: 0.040556274354457855
epoch: 0 batch: 1200 loss: 0.08253474533557892
epoch: 0 batch: 1800 loss: 0.3647049069404602
epoch: 0 batch: 2400 loss: 0.018250251188874245
epoch: 0 batch: 3000 loss: 0.008067040704190731
epoch: 0 batch: 3600 loss: 0.001166942878626287
epoch: 0 batch: 4200 loss: 0.5255253911018372
epoch: 0 batch: 4800 loss: 0.03260819613933563
epoch: 0 batch: 5400 loss: 0.007468158844858408
epoch: 0 batch: 6000 loss: 0.03889675810933113
epoch: 1 batch: 600 loss: 0.032828204333782196
epoch: 1 batch: 1200 loss: 0.04554177075624466
epoch: 1 batch: 1800 loss: 0.005784796085208654
epoch: 1 batch: 2400 loss: 0.02235613949596882
epoch: 1 batch: 3000 loss: 0.21643038094043732
epoch: 1 batch: 3600 loss: 0.00501451687887311
epoch: 1 batch: 4200 loss: 0.00045869071618653834
epoch: 1 batch: 4800 loss: 0.0019295118981972337
epoch: 1 batch: 5400 loss: 0.0008596166153438389
epoch: 1 batch: 6000 loss: 0.08304359018802643
epoch: 2 batch: 600 loss: 0.0006373372743837535
epoch: 2 batch: 1200 loss: 0.0015393418725579977
epoch: 2 batch: 1800 loss: 0.0012801657430827618
epoch: 2 batch: 2400 loss: 0.001396776526235044
epoch: 2 batch: 3000 loss: 0.3044474124908447
epoch: 2 batch: 3600 loss: 0.014451900497078896
epoch: 2 batch: 4200 loss: 0.021982822567224503
epoch: 2 batch: 4800 loss: 0.0007802899926900864
epoch: 2 batch: 5400 loss: 0.0016833205008879304
epoch: 2 batch: 6000 loss: 0.0002076365490211174
epoch: 3 batch: 600 loss: 0.0007947428966872394
epoch: 3 batch: 1200 loss: 0.002038671402260661
epoch: 3 batch: 1800 loss: 0.0004689941997639835
epoch: 3 batch: 2400 loss: 0.00021815943182446063
epoch: 3 batch: 3000 loss: 0.031423646956682205
epoch: 3 batch: 3600 loss: 0.0073494575917720795
epoch: 3 batch: 4200 loss: 0.0006103587802499533
epoch: 3 batch: 4800 loss: 0.13828447461128235
epoch: 3 batch: 5400 loss: 0.0007458419422619045
epoch: 3 batch: 6000 loss: 0.02092968113720417
epoch: 4 batch: 600 loss: 0.0009378452086821198
epoch: 4 batch: 1200 loss: 0.19402171671390533
epoch: 4 batch: 1800 loss: 0.0006758190575055778
epoch: 4 batch: 2400 loss: 0.00019682350102812052
epoch: 4 batch: 3000 loss: 0.005403806921094656
epoch: 4 batch: 3600 loss: 0.0005835095071233809
epoch: 4 batch: 4200 loss: 0.0011737591121345758
epoch: 4 batch: 4800 loss: 0.0018565601203590631
epoch: 4 batch: 5400 loss: 0.0002484446158632636
epoch: 4 batch: 6000 loss: 0.00652279332280159
Training took 4.294626307487488 minutes
```

```
plt.plot(train_losses, label='training loss')
plt.plot(test_losses, label='validation loss')
plt.title('Loss at the end of each epoch')
plt.legend()
plt.show()
```



```
plt.plot([t/600 for t in train_correct], label='training accuracy')
plt.plot([t/100 for t in test_correct], label='validation accuracy')
plt.title('Accuracy at the end of each epoch')
plt.legend();
plt.show()

# Extract the data all at once, not in batches
test_load_all = DataLoader(test_data, batch_size=10000, shuffle=False)


with torch.no_grad():
    correct = 0
    for X_test, y_test in test_load_all:
        y_val = model(X_test) # we don't flatten the data this time
        predicted = torch.max(y_val,1)[1]
        correct += (predicted == y_test).sum()

# print a row of values for reference
np.set_printoptions(formatter=dict(int=lambda x: f'{x:4}'))
print(np.arange(10).reshape(1,10))
print()
```

```
# print the confusion matrix
print(confusion_matrix(predicted.view(-1), y_test.view(-1)))
```

```
[[ 0  1  2  3  4  5  6  7  8  9]]
[[ 977  3  1  0  0  2  4  1  4  0]
 [  0 1130  2  0  0  0  1  4  0  3]
 [  0  1 1022  0  0  0  0  4  2  0]
 [  0  0  3 1007  0  7  0  2  2  2]
 [  0  0  1  0 971  0  1  0  0  5]
 [  0  0  0  1  0 879  7  0  1  5]
 [  1  1  0  0  4  2 944  0  1  0]
 [  1  0  3  0  0  0  0 1013  1  2]
 [  1  0  0  2  1  1  1  1 962  8]
 [  0  0  0  0  6  1  0  3  1 984]]
```

```
# single image for test
plt.imshow(test_data[2019][0].reshape(28,28))
plt.show()
```



```
model.eval()
with torch.no_grad():
    new_prediction = model(test_data[2019][0].view(1,1,28,28))
```

```
new_prediction.argmax()
```

```
tensor(9)
```

```
torch.save(model.state_dict(), 'yogesh212222110055.pt')
```

```
new_model = ConvolutionalNetwork() # Replace Model with ConvolutionalNetwork
new_model.load_state_dict(torch.load('yogesh212222110055.pt'))
new_model.eval()
```

```
↔ ConvolutionalNetwork(  
  (conv1): Conv2d(1, 6, kernel_size=(3, 3), stride=(1, 1))  
  (conv2): Conv2d(6, 16, kernel_size=(3, 3), stride=(1, 1))  
  (fc1): Linear(in_features=400, out_features=120, bias=True)  
  (fc2): Linear(in_features=120, out_features=84, bias=True)  
  (fc3): Linear(in_features=84, out_features=10, bias=True)  
)
```